

CS 461

Artificial Intelligence

Dr. Hashim Yasin

Learning Agent

Learning Agent

- ▶ **Turing** – instead of actually programming **intelligent machines** by hand, build **learning machines** and then teach them
- ▶ All agents can improve their performance through **learning**
- ▶ Learning also allows the agent
 - to operate in initially unknown environments
 - to become more competent than its initial knowledge

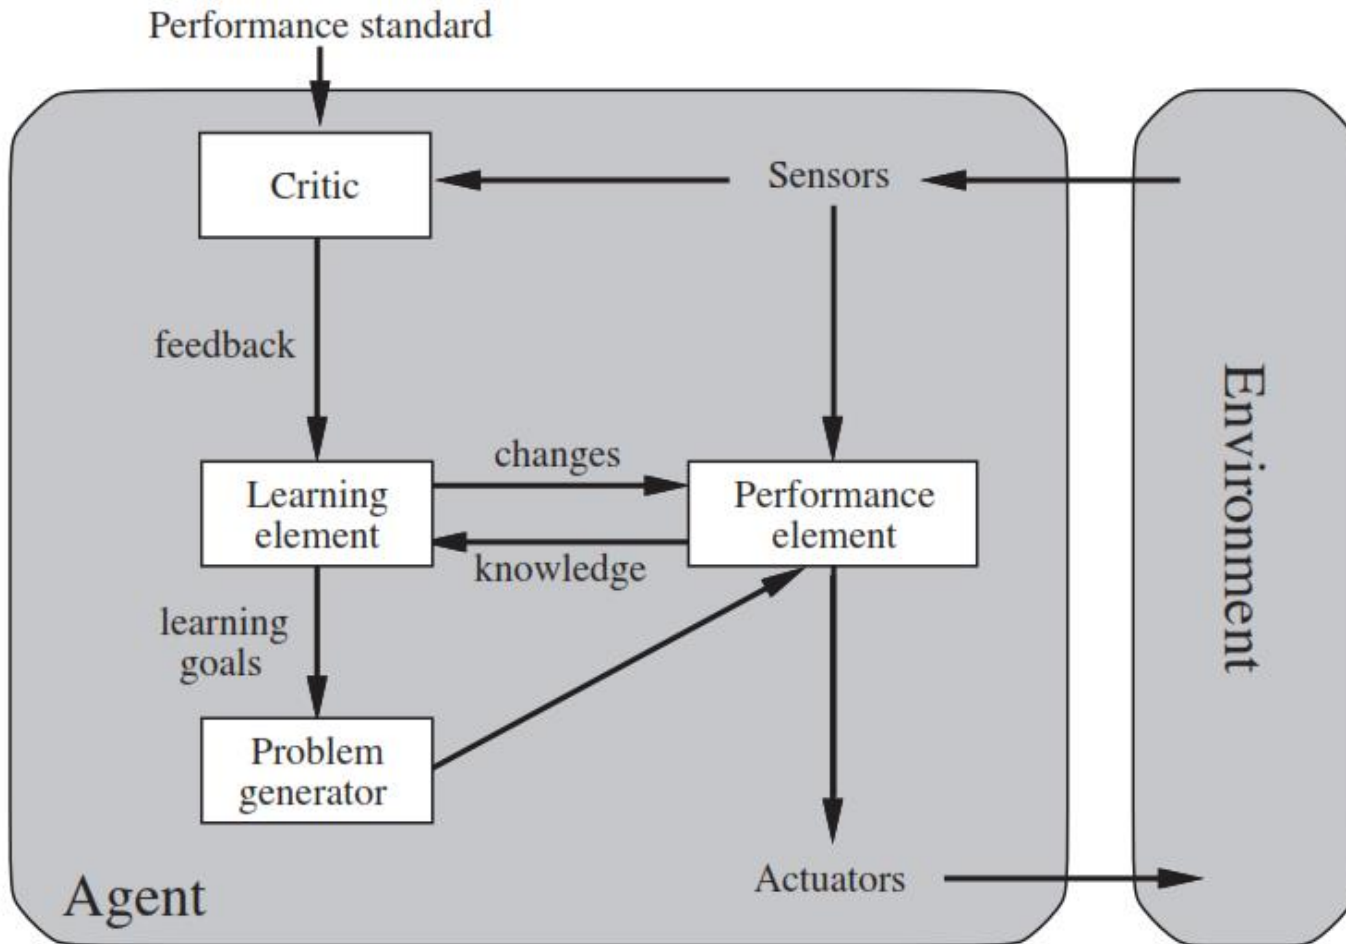
Learning Agent

- ▶ A learning agent can be divided into four conceptual components:
 1. The **learning element**, which is responsible for *making improvements*,
 2. The **performance element**, which is responsible for *selecting external actions*.
 3. The learning element uses feedback from the **critic**
 - how the agent is doing and
 - determines how the performance element should be modified *to do better* in the future.

Learning Agent

4. The last component of the learning agent is the **problem generator**.
 - It is responsible for suggesting actions that will lead to *new and informative experiences*.
 - The problem generator's job is to suggest these exploratory actions.

Learning Agent



State Representations

State Representations

State:

- ▶ A state is the representation of the **physical configuration**. It may be considered as all relevant aspects of the problem.
- ▶ Three ways to represent states and the transitions between them.

Atomic representation:

- ▶ A state (such as **B** or **C**) is a black box with no internal structure;
 - The algorithms underlying **search, game-playing, Hidden Markov models, and Markov decision processes** etc.

State Representations


Factored representation:

- ▶ A state consists of **a vector of attribute values**; values can be Boolean, real valued, or one of a fixed set of symbols.
 - Constraint Satisfaction Problems, Propositional Logic, Planning, Bayesian Networks


Structured representation:

- ▶ A state includes **objects**, each of which may have attributes of its own as well as relationships to other objects
 - First-order Logic, Knowledge-based Learning


State Representations



(a) Atomic



(b) Factored



(b) Structured

Problem Solving Agents



Problem Solving Agents

- ▶ **Problem formulation** is the process of deciding what actions and states to consider, given a goal
- ▶ A **goal-based agent** is also called a **problem-solving agent**.
- ▶ The problem-solving agents use *atomic representation*.
- ▶ Goal-based agents that use more advanced **factored** or **structured** representations are usually called **planning agents**.

Searching

- ▶ *How can an agent find a sequence of actions that achieves its goal when no single action will do?*
- ▶ One general approach to problem solving is
 - **to reduce the problem to be solved to one of searching a tree or graph.**
- ▶ Search plays a key role in many parts of AI.

Trees & Graphs

Tree:

- ▶ A tree is made up of **nodes and links (circles and lines)** connected so that there are NO loops (cycles).
- ▶ *Nodes are sometimes referred to as vertices* and **links as edges** (this is more common in graphs).
- ▶ A tree has a **root node** (where the tree "starts").
- ▶ Every node except the root has a single parent (aka direct **ancestor**).
- ▶ Each node except the terminal (**leaf nodes**) has one or more children (**direct descendants**).

Trees & Graphs

Graph:

- ▶ A **graph** is also a set of nodes connected by links but where **loops are allowed** and a node can have multiple parents.
- ▶ We have two kinds of graphs to deal with: Directed graphs and undirected graphs
- ▶ **Directed graphs**, where the links have direction (for example: one-way streets).
- ▶ **Undirected graphs**, where the links go both ways.

A Problem Definition

- ▶ A **problem** can be defined formally by five components:
 - ❑ Initial State
 - ❑ Actions
 - ❑ Transition Model
 - ❑ Goal Test
 - ❑ Path Cost

A Problem Definition

- ▶ **State**: A state is the representation of the **physical configuration**. It may be considered as all relevant aspects of the problem.
- ▶ **Actions**: Actions are operators:
 - ▶ **Actions are deterministic**, when we know exactly the state after an action is performed.
 - ▶ **Actions are discrete**, when we don't have to represent what happens while the action is happening.
 - For example, we assume that a flight gets us to the scheduled destination and that what happens during the flight does not matter.

A Problem Definition

- ▶ **Goal Test:** It determines whether a given state is a goal state.
- ▶ In general, we need a test for the goal, *not just one specific goal state*.
- ▶ For example,
 - We might be interested in any city in Germany rather than specifically Frankfurt.
 - When proving a theorem, all we care is about knowing one fact in our current database of facts.
 - Any final set of facts that contains the desired fact is a proof.
- ▶ **Path Cost:** Some penalty has been allocated.

A Problem Definition


Transition Model:

- ▶ A description of what does each action do;
- ▶ The **transition model**, specified by a function $RESULT(s, a)$ that returns the state that results from doing action a in state s .
- ▶ We also use the term **successor** to refer to any state reachable from a given state by a single action

$$RESULT(In(FSD), Go(LHR)) = In(LHR)$$


Toy Problems

Example: vacuum world



- ▶ **States:** The state is determined by both the **agent location (2)** and the **dirt locations (2^2)**: two locations with or without dirt: **$2 \times 2^2 = 8$ states**.
 - larger environment with **n** locations has **$n \times 2^n$ states**
- ▶ **Initial state:** Any state can be initial

Example: vacuum world



- ▶ **Actions:** {*Left*, *Right*, *Suck*}
- ▶ **Transition model:** The actions have their expected effects, except that moving *Left*, moving *Right*, and *Sucking* in a clean square have no effect.
- ▶ **Goal test:** Check whether all squares are clean.
- ▶ **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

Example: 8-puzzle Problem

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- ▶ **States:** Integer location of each tile
 - A state description specifies the location of each of the eight tiles and the blank in one of the nine squares.
- ▶ **Initial state:** Any state can be initial
- ▶ **Actions:** {*Left, Right, Up, Down*}

Example: 8-puzzle Problem

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State


- ▶ **Transition model:** Given a state and action, this returns the resulting state;
 - for example, if we apply *Left* to the start state in Figure, the resulting state has the 5 and the blank switched.
- ▶ **Goal test:** Check whether goal configuration is reached.
- ▶ **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.

Example: 8-puzzle Problem


- ▶ The **8-puzzle** has 181, 440 reachable states and is easily solved.
- ▶ The **15-puzzle** (on a 4×4 board) has around 1.3 trillion states, and random instances can be solved optimally in a few milliseconds by the best search algorithms.
- ▶ The **24-puzzle** (on a 5 × 5 board) has around 10^{25} states, and random instances take several hours to solve optimally.

Example: 8-queen Problem

- ▶ Place 8 queens in a chessboard so that no two queens are in the same row, column, or diagonal.



A solution



Not a solution


Example: 8-queen Problem

- ▶ **States:** any arrangement of 0 to 8 queens on the board
- ▶ **Initial state:** No queen on the board
- ▶ **Transition model:** Returns the board with a queen added to the specified square.
- ▶ **Actions:** add a queen in any square
- ▶ **Goal test:** 8 queens on the board, none attacked
- ▶ **Path cost:** none

64^8 states with 8 queens

Real Life Problems

Example: Robot Assembly



- ▶ **States:** Real-valued coordinates of robot joint angles; parts of the object to be assembled.
- ▶ **Initial state:** Any arm position and object configuration.
- ▶ **Actions:** Continuous motion of robot joints
- ▶ **Goal test:** Complete assembly
- ▶ **Path cost:** Time to execute

Example: Route Finding Problem

- ▶ **States:** each state is represented by **a location** (e.g., An airport) and the **current time**
- ▶ **Initial state:** specified by the problem
- ▶ **Actions:** *Take any flight from the current location*, in any seat class, leaving after the current time, leaving enough time for within-airport transfer if needed.

Example: Route Finding Problem

- ▶ **Transition model:** The state resulting from taking a flight will have the *flight's destination* as the **current location** and the *flight's arrival time* as the **current time**.
- ▶ **Goal test:** are we at the destination by some pre-specified time
- ▶ **Path cost:** monetary cost, waiting time, flight time, customs and immigration procedures, seat quality, time of day, type of airplane, frequent-flyer mileage awards, etc.

Searching

Uninformed search

- ▶ No information about the number of steps
- ▶ No path cost taken into account from current state to the goal
- ▶ Search the state space **blindly**

Informed search, or heuristic search

- ▶ A cleverer strategy is used that searches toward the goal.
- ▶ Based on the information from the current state.

Searching

Any-path searches and optimal searches.

- ▶ **Any-path searches** will just settle for finding some solution,
- ▶ **Optimal searches** are looking for the **best possible path**.

Evaluation of a Search Strategy

- ▶ **Completeness:** is the strategy guaranteed to find a solution when there is one?
- ▶ **Optimality:** does the strategy find the highest-quality (least-cost) solution when there are several different solutions?
- ▶ **Time complexity:** how long does it take to find a solution?
- ▶ **Space complexity:** how much memory is needed to perform the search?

Evaluation of a Search Strategy

Complexity may be expressed in

- ▶ **b**, branching factor, maximum number of successors of any node
- ▶ **d**, the depth of the shallowest goal node. (number of steps along the path from the root), depth of the least-cost solution
- ▶ **m**, the maximum length of any path in the state space

Time and Space is measured in

- ▶ number of nodes generated during the search
- ▶ maximum number of nodes stored in memory

Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - Chapter 2 & 3.

