

CS 461

Artificial Intelligence

Dr. Hashim Yasin

Beyond Classical Search

Beyond Classical Search

- ▶ We have addressed a single category of problems: **observable, deterministic, known environments**
 - where the solution is a sequence of actions
- ▶ The search algorithms that we have seen so far are designed to explore search spaces **systematically**.
 - When a **goal is found**, the **path** to that goal also constitutes a **solution** to the problem.


What happens when these assumptions are *relaxed*?

Beyond Classical Search

- ▶ In many problems, however, the **path to the goal is irrelevant.**

For example:

- ▶ 8-queens problem
 - *what matter, is the final configuration of queens, not the order.*
- ▶ The factory-floor layout problem
- ▶ The vehicle routing problem



Local Search Algorithms

Local Search Algorithms

- ▶ Operate using a **single current node** and generally **move only to neighbors** of that node.
 - *No concern with paths* followed by the search
 - They are *NOT systematic*
- ▶ Local search algorithms have **two key advantages**:
 - they use very *little memory*
 - they can often find reasonable *solutions in large* or *infinite (continuous)* state spaces

Local Search Algorithms

- ▶ The local search algorithms are useful for solving pure **optimization problems**,

Optimization problems:


- ▶ the aim is **to find the best state** according to an *objective function*.
- ▶ The *standard search methods do not fit well* with optimization problems.

Local Search Algorithms

State-space landscape

A landscape has both

- ▶ **“location”**
 - defined by the state
- ▶ **“elevation”**
 - defined by the value of the **heuristic cost function** or **objective function**.




Local Search Algorithms

State-space landscape

- ▶ If elevation corresponds to cost, then the aim is to find the lowest valley—a **global minimum**;
- ▶ If elevation corresponds to an **objective function**, then the aim is to find the highest peak—a **global maximum**.

Local Search Algorithms

State-space landscape




A **one-dimensional state-space landscape** in which *elevation corresponds to the objective function*. The aim is to find the **global maximum**.

Local Search Algorithms


- ▶ A **complete local search** algorithm always finds a goal if one exists;
- ▶ An **optimal** algorithm always finds a global minimum/maximum.
- ▶ Local search algorithms typically use a **complete-state formulation**,
 - In 8-queen problem, where *each state has 8 queens* on the board, one per column.

Example: n -queens Problems


- ▶ Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- ▶ Move a queen to reduce number of conflicts
- ▶ Heuristic h : number of 'attacks'



$h = 5$



$h = 2$






$h = 0$

Example: n -queens Problems

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts
- Heuristic h : number of 'attacks'

If queen is placed here, there is only one attack then,



Hill-climbing Search

Hill-climbing Search

- ▶ It is simply a **loop** that continuously moves in the direction of increasing value—that is, uphill.
- ▶ It terminates when it reaches a “peak” where no neighbor has a higher value.
 - *does not maintain a search tree*
 - need only
 - *record the state* and
 - the value of the *objective function*.
- ▶ Hill climbing only looks towards the **immediate neighbors** of the current state.

Hill-climbing Search

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*


if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*

- ▶ At *each step the current node is replaced by the best neighbor*; the neighbor with the highest VALUE,
- ▶ If a **heuristic cost estimate h** is used, we would find the neighbor with the ***lowest h*** .

Hill-climbing Search

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♙	13	16	13	16
♙	14	17	15	♙	14	16	16
17	♙	16	18	15	♙	15	♙
18	14	♙	15	15	14	♙	16
14	14	13	17	12	14	12	18



- (a) An 8-queens state with heuristic cost estimate $h = 17$, showing the value of h for each possible successor. (b) A local minimum: the state has $h = 1$ but every successor has a higher cost.

Hill-climbing Search

- ▶ The hill climbing often gets stuck for the following reasons:


Local maxima:

- ▶ A local maximum is a peak that is higher than each of its neighboring states but lower than the global maximum.

Hill-climbing Search

Ridges:

- ▶ A sequence of local maxima
- ▶ *Its difficult for greedy algorithms to navigate*




Hill-climbing Search

Plateaux:


- ▶ An area of the state space where the **evaluation function is flat**.
- ▶ It can be a **flat local maximum**, from which no uphill exit exists,

Shoulder:

- ▶ from which progress is possible



Hill-climbing Search



Reading Material

- ▶ **Artificial Intelligence, A Modern Approach**
Stuart J. Russell and Peter Norvig
 - **Chapter 3 & 4.**

