

CS 461

ARTIFICIAL INTELLIGENCE

Lecture # 03

March 10, 2021

SPRING 2021

FAST – NUCES, CFD Campus

Dr. Rabia Maqsood
rabia.maqsood@nu.edu.pk

Today's Topics

- Problem solving agents
 - *Introduction*
 - *Understanding search problems*
- State space graph
- Search strategies
 - *Uninformed search algorithms*
 - Depth First Search
 - Breadth First Search

Search

- Finding the sequence of actions to achieve some goal(s) given a start state is called **search**, and the agent which does this is called the *problem-solving agent*.

Search

- Often, we are not given an algorithm to solve a problem, but only a specification of what a solution is, and

*we have to **search** for a solution.*

- A typical approach:
 - *Enumerate a set of potential partial solutions*
 - *Check to see if they are solutions or could lead to one*

A Simple Search Agent

Deterministic, goal-driven agent


- Agent is in a **start state**
- Agent is given a **goal** (subset of possible states)
- Environment changes only when the agent acts
- Agent perfectly knows:
 - **actions** that can be applied in any given state
 - the **state** it is going to end up in when an action is applied in each state
- The sequence of actions (and appropriate ordering) taking the agent from the start state to a goal state is the **solution**

Definition of a search problem

- **Initial state(s)**
- Set of **actions** (operators) available to the agent
- An **action function (or transition model)** that, given a state and an action, returns a new state
- **Goal state(s)**
- By combining initial state, actions and transition model we get, **state space**
 - *State space is a set of states that will be searched for a path from initial state to goal, given the available actions*
 - **states are nodes** and **actions are links** between them.
 - *Not necessarily given explicitly (state space might be infinite)*
- **Path Cost** (we ignore this for now)

Example: vacuum world

- States
 - *Two rooms: $r1$, $r2$*
 - *Each room can be either dirty or not*
 - *Vacuum agent can be in either in $r1$ or $r2$*
 - *How many total states?* **8**




Possible start state




Possible goal state

Example: vacuum world




- **States** – one of the eight states in the picture
- **Actions** – *left, right, suck*
- **Possible Goal** – no dirt

Search Space



- **Actions** – left, right, suck
 - Successor states in the graph describe the effect of each action applied to a given state
- **Possible Goal** – no dirt

Eight Puzzle



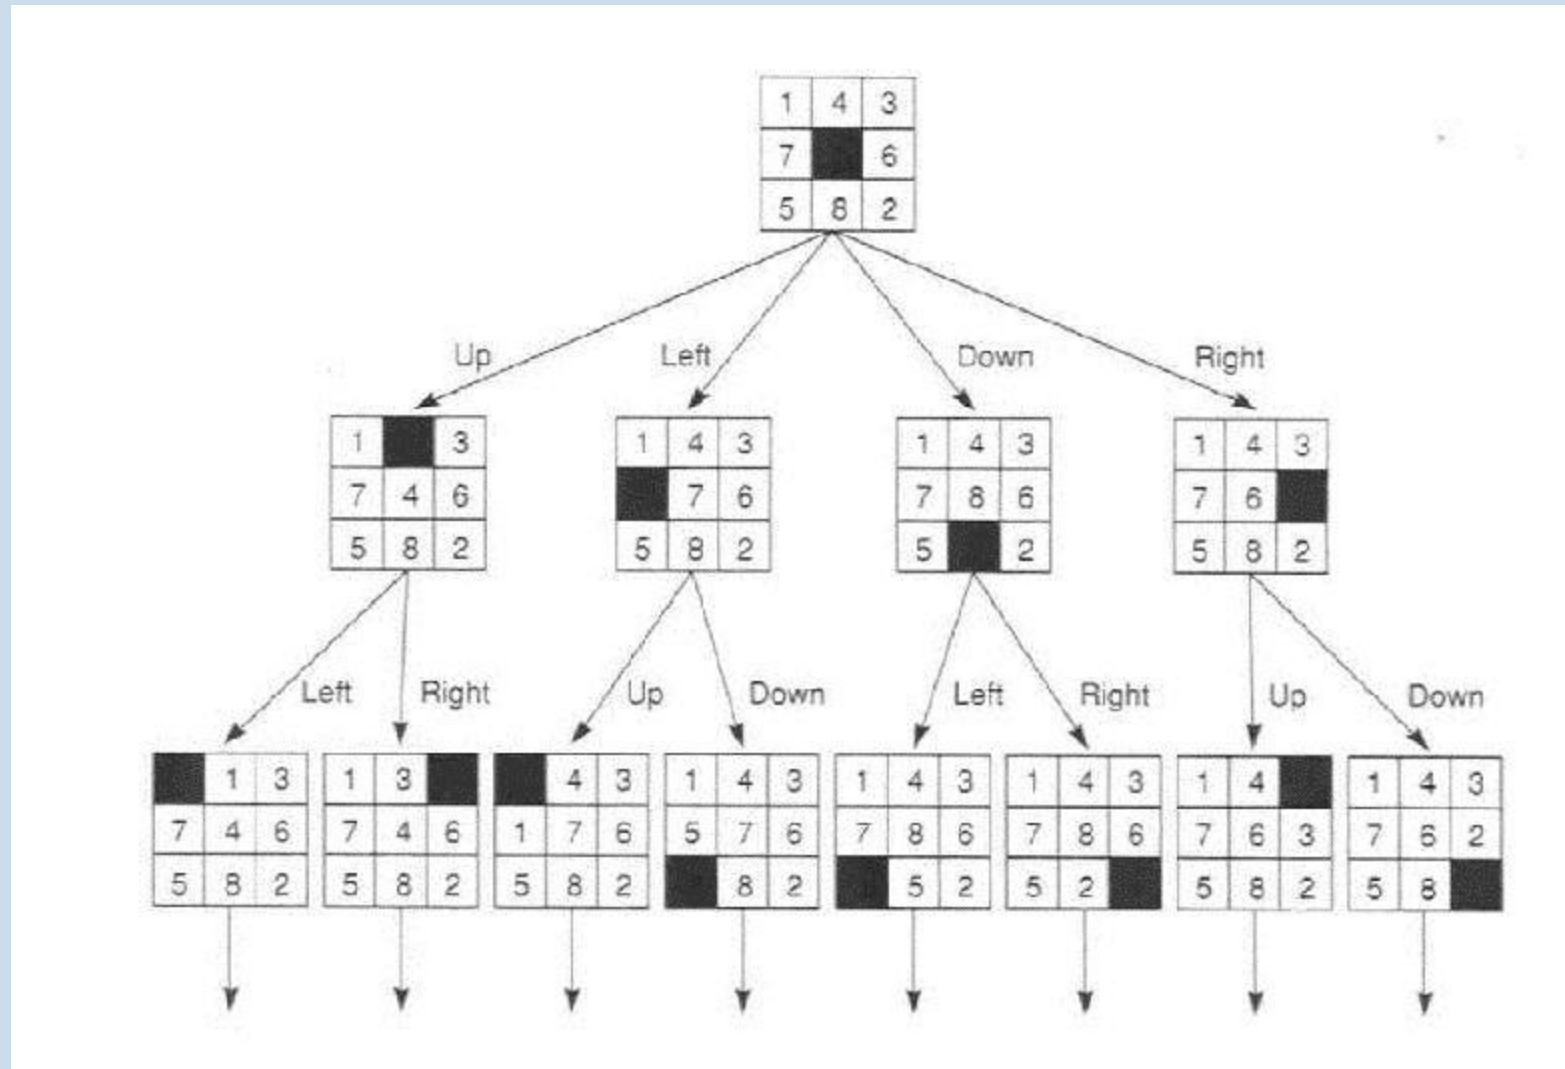
States: each state specifies which number/blank occupies each of the 9 tiles

HOW MANY STATES ? **9!**

Actions: blank moves left, right, up down


Goal: configuration with numbers in right sequence

Search space for Eight Puzzle



Search problem representation: “Just” a Graph

- States: nodes
- Actions: edges
- Path: sequence of edges



How can we find a solution?

- How can we find a sequence of actions and their appropriate ordering that led to the goal?
- Need smart ways to **search** the **state space graph**

Graph search: basic idea

- Given a graph, start node, and goal nodes, incrementally explore paths from the start node.
- Maintain a *frontier* of paths from the start node that have been explored.
- As search proceeds, the *frontier* expands into the unexplored nodes until a goal node is encountered.

Graph search: basic idea

Input:

- a graph
- a set of start nodes
- Boolean procedure `goal(n)` testing if `n` is a goal node

`frontier` := [`<s>`: `s` is a start node];

While `frontier` is not empty:

 select and remove path

`<n0, ..., nk>` from

`frontier`;

 If `goal(nk)`

 return `<n0, ..., nk>`;


 For every neighbor `n` of `nk`,

 add `<n0, ..., nk, n>` to `frontier`;

end

CS 461 - SPRING 2021

The way in which the frontier is expanded defines the **search strategy**



Comparing Search Strategies

- Algorithms are evaluated along the following four dimensions:

1. *Completeness*
2. *Optimality*
3. *Time complexity*
4. *Space complexity*

Comparing Search Strategies

1. Completeness

- A search algorithm is **complete** if whenever there is at least one solution, the algorithm **is guaranteed to find it** within a finite amount of time.

2. Optimality

- A search algorithm is **optimal** if when it finds a solution, it is **the best one**

Comparing Search Strategies

3. Time complexity

- The **time complexity** of a search algorithm is the **worst-case** amount of time it will take to run, expressed in terms of:
 - **b** – maximum branching factor of the search tree
 - **d** – depth of the least-cost solution
 - **m** – maximum depth of the state space

4. Space complexity

- The **space complexity** of a search algorithm is the **worst-case** amount of memory that the algorithm will use (i.e., the maximum number of nodes on the frontier), also expressed in terms of **b** , **d** and **m** .

UNINFORMED SEARCH STRATEGIES

A.K.A “Blind Search”


Uninformed Search Strategies

- *Uninformed* strategies use only the information available in the problem definition (i.e., no additional information is available)
 - *Breadth-first search*
 - *Uniform cost search*
 - *Depth-first search*
 - *Depth limited search*
 - *Iterative deepening search*
 - *Bidirectional search*

Depth-First Search

In DFS, the frontier is a
last-in-first-out (stack)


- Basic idea
 - explores each path on the frontier until its end (or until a goal is found) before considering any other path.
- Example: run the DFS for the given graph.
 - Take 0 as the start state & 5 as the goal state
 - Use lexicographic order to put states into the frontier (i.e., a stack in this case)



Depth-First Search

In DFS, the frontier is a
last-in-first-out (stack)

- Basic idea
 - explores each path on the frontier until its end (or until a goal is found) before considering any other path.
- Implementation
- Properties
 - Complete?
 - Optimal?
 - Time?
 - Space?




Depth-First Search

- Complete? No
 - *fails in infinite-depth spaces, spaces with loops*
 - *Complete in finite spaces*
- Optimal? No
- Time? $O(b^m)$
 - *terrible if m is much larger than d*
 - *But if solutions are dense, may be much faster than breadth-first*
- Space? $O(bm)$, i.e., linear space!
- Preventing loop paths?

Breadth-First Search

In BFS, the frontier is a
first-in-first-out (queue)

- Basic idea
 - explores all paths of length k on the frontier, before looking at path of length $k + 1$
- Implementation
- Properties
 - Complete?
 - Optimal?
 - Time?
 - Space?



Breadth-First Search

- Complete? Yes (if b is finite)
- Optimal? Yes (if cost = 1 per step); not optimal in general
- Time? $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$ (i.e., exponential in d)
- Space? $O(b^d)$ (keeps every node in memory)
- **Space** is the bigger problem (more than time).

BFS vs. DFS: which one to use?

- The search graph has cycles or is infinite

BFS

- We need the shortest path to a solution

BFS

- There are only solutions at great depth

DFS

- There are some solutions at shallow depth and others deeper

BFS

- No way the search graph will fit into memory

DFS

Reading Material

- Russell & Norvig: Chapter # 3
- David Poole: Chapter # 3